CS532 Web Science: Assignment 3

Finished on February 18, 2016

Dr. Michael L. Nelson

Naina Sai Tipparti

ntippart@cs.odu.edu

Contents

Listings

List of Figures

Problem 1

Question

Download the 1000 URIs from assignment #2. "curl", "wget", or "lynx" are all good candidate programs to use. We want just the raw HTML, not the images, stylesheets, etc.

from the command line:

% curl http://www.cnn.com/ > www.cnn.com

% wget -O www.cnn.com http://www.cnn.com/

% lynx -source http://www.cnn.com/ > www.cnn.com

"www.cnn.com" is just an example output file name, keep in mind that the shell will not like some of the characters that can occur in URIs (e.g., "?", "&"). You might want to hash the URIs, like:

% echo -n "http://www.cs.odu.edu/show_features.shtml?72" |md5
41d5f125d13b4bb554e6e31b6b591eeb

("md5sum" on some machines; note the "-n" in echo - this removes the trailing newline.)

Now use a tool to remove (most) of the HTML markup. "lynx" will do a fair job:

% lynx -dump -force_html www.cnn.com > www.cnn.com.processed
Use another (better) tool if you know of one. Keep both files for each URI (i.e., raw HTML and processed).

Answer

Using the python script in Listing 1, 1000 unique URIs were dereferenced and their raw contents were stored in the `html/raw/` folder as a file with the filename as the md5-hashed URI. These were then stripped of all html elements and their processed contents were stored in the `html/processed/` folder as the same md5-hashed filename. For reference, the URIs were written as the first line of each of their content files.

```python
#! /usr/bin/python

import requests
import concurrent.futures
import md5
from bs4 import BeautifulSoup
import pickle

def convert(uri):
```

```
            return md5.new(uri).hexdigest()
11
    def get_html(uri):
            print('Getting {}'.format(uri))
            response = requests.get(uri)
            return response.url, response.status_code, response.content
16
    if __name__ == '__main__':
            with open('links') as infile:
                    uris = [uri.rstrip('\n') for uri in infile]

21          with concurrent.futures.ThreadPoolExecutor(max_workers=8) as executor:
                    uri_futures = [executor.submit(get_html, uri) for uri in uris]
                    for future in concurrent.futures.as_completed(uri_futures):
                            try:
                                    uri, status_code, content = future.result()
26                          except Exception as exc:
                                    print('{} generated an exception: {}'.format(uri, exc))
                                    continue
                            if status_code == 200:
                                    hashed_uri = convert(uri)
31                                  print('Writing {} as {}'.format(uri, hashed_uri))
                                    try:
                                            with open('html/raw/' + hashed_uri, 'w') as outfile:
                                                    outfile.write(uri + '\n')
                                                    outfile.write(content)
36                                          with open('html/processed/' + hashed_uri + '.processed.txt', 'w') as
                                                    outfile:
                                                    outfile.write(uri + '\n')
                                                    outfile.write(BeautifulSoup(content).get_text().encode('utf8'))
                                    except Exception as e:
                                            print '**** ERROR **** --- ' + uri
41                                          print e
                            else:
                                    print('Not writing {}, bad status code: {}'.format(uri, status_code))
```

Listing 1: get_html.py

Problem 2

Question

Choose a query term (e.g., "shadow") that is not a stop word (see week 5 slides) and not HTML markup from step 1 (e.g., "http") that matches at least 10 documents (hint: use "grep" on the processed files). If the term is present in more than 10 documents, choose any 10 from your list. (If you do not end up with a list of 10 URIs, you've done something wrong).

As per the example in the week 5 slides, compute TFIDF values for the term in each of the 10 documents and create a table with the TF, IDF, and TFIDF values, as well as the corresponding URIs. The URIs will be ranked in decreasing order by TFIDF values. For example:

Table 1. 10 Hits for the term "shadow", ranked by TFIDF.

TFIDF TF IDF URI
0.150 0.014 10.680 http://foo.com/
0.044 0.008 5.510 http://bar.com/

You can use Google or Bing for the DF estimation. To count the number of words in the processed document (i.e., the deonminator for TF), you can use "wc":
    % wc -w www.cnn.com.processed
2370 www.cnn.com.processed

It won't be completely accurate, but it will be probably be consistently inaccurate across all files. You can use more accurate methods if you'd like.

Don't forget the log base 2 for IDF, and mind your significant digits!

Answer

First, the function `count_terms` was used to count the term frequency for the given term "shadow" in all documents.

```
def count_terms(term, file_list=os.listdir('html/processed')):
    for filename in file_list:
15      with open('html/processed/' + filename + '.processed.txt') as infile:
            uri = infile.readline().strip()
            text = infile.read()
            count = text.count(term)
            if count > 0:
20              print('{} {}'.format(count, uri))
                return count, uri
```

Listing 2: count_terms function

Ten of the results were chosen at random and stored in the `uri_counts` file. In order to easily identify which file corresponds to which URI, since the filename is the non-reversible md5-hashed URI string, a mapping from URI to filename was created using the functions in Listing 3 and serialized in the `uri_map` file using the `pickle` library.

```
def get_uri(uri):
25  for filename in os.listdir('html/processed/'):
        with open('html/processed/' + filename + '.processed.txt') as infile:
            if uri in infile.readline():
                return uri, filename
    return None, None
30
def get_uris():
    uri_file = {}
    for uri in open('links').read().split('\n'):
        uri, filename = get_uri(uri)
35      if not uri:
            continue
        uri_file[uri] = filename
```

Listing 3: get_uris functions

Reading from the file was done with the line in Listing 4. This loaded the serialized URI to filename map for future use.

```
10  map_file = open('uri_map', 'rb')
```

Listing 4: Loading the uri map

To proceed with processing each of the files to find Term Frequency (TF), Inverse Document Frequency (IDF) and the product of the two (TFIDF), each URI's corresponding file was found using the `get_filename` function found in Listing 5.

```
40  def get_filename(uri):
        if uri_map.has_key(uri):
            return uri_map[uri]
```

Listing 5: Getting filename from URI

Then, they were stripped of HTML tags using the `strip_html` function in Listing 6.

```python
45  def strip_html(filename):
        if not filename:
            print 'invalid filename'
            return
        with open('html/processed/' + filename + '.processed.txt') as infile:
50          # To remove URI in first line
            infile.readline()
            # Removing all punctuation
            strs = infile.read()
            r = re.compile(r'[{}]'.format(punctuation))
55          content = r.sub(' ', strs)
```

Listing 6: Stripping HTML tags from content

And finally the frequencies were calculated for each URI using the functions in Listing 7.

```python
    def get_tf(content, term):
        return float(content.count(term)) / float(len(content.split()))
60
    def get_idf(term):
        present = set()
        absent = set()
        for uri, filename in uri_map.iteritems():
65          content = strip_html(filename)
            if not content:
                continue
            if term in content:
                present.add(uri)
70          else:
                absent.add(uri)
        return math.log(float(len(absent)) / float(len(present)), 2)

    def process_uri(uri, term):
75      tf = get_tf(strip_html(get_filename(uri)), term)
        tfidf = tf * idf
        return tf, tfidf
```

Listing 7: Calculating TF, IDF & TFIDF

These frequencies were then written to the `uri_frequencies` file using the code in 8.

```python
            term = sys.argv[2]
            with open('uri_counts') as infile:
                uris = uris = [line.split()[1] for line in infile.read().split('\n')]
110             with open('uri_frequencies', 'w') as outfile:
                    outfile.write('{:<7} {:<7} {:<7} {:<7}\n'.format('TFIDF', 'TF', 'IDF', 'URI'
                        ))
                    for uri in uris:
                        tf, tfidf = process_uri(uri, term)
```

Listing 8: Writing results to uri_frequencies file

And the results can be seen in Listing 9:

```
TFIDF    TF       IDF        URI
_____    __       ___        ___
0.0220   0.0065   3.3825     http://news.google.com/
0.0159   0.0047   3.3825     http://www.easkme.com/2014/07/mail-merge-in-gmail.html#.VB8GnkIk6rA.
    facebook
0.0113   0.0033   3.3825     http://btc-news-bot.tumblr.com/post/98066358831/we-need-to-do-a-
    better-job-explaining-bitcoin-in-ways#_=_
0.0109   0.0032   3.3825     http://musicisthedrug-revolution.tumblr.com/post/98067184737/
    description-you-know-what-time-it-is-world-cup#_=_
0.0099   0.0029   3.3825     http://www.ebay.com/itm/4-5-Android-Smartphone-Dual-Sim-Dual-Core-
    Unlocked-WIFI-3G-GSM-Cell-phone-AT-T-/271610837947?pt=Cell_Phones&hash=item3f3d447bbb
0.0092   0.0027   3.3825     http://www.blogdeizquierda.com/2014/09/mundo-bitcoin-banqueros-de-eu
    -ven.html?utm_source=twitterfeed&utm_medium=twitter
0.0064   0.0019   3.3825     http://www.valuewalk.com/2014/09/best-apps-apple-iphone-6/
0.0064   0.0019   3.3825     http://www.datelinemovies.com/2014/07/bloopers-for-season-4-game-of-
    thrones.html#sthash.yYuV0eDx.uxfs
0.0027   0.0008   3.3825     http://abusidiqu.com/its-all-scripted-ebola-virus-is-a-biological-
    weapon-from-the-us-read-this-shocking-report/
0.0019   0.0006   3.3825     http://rss-now.blogspot.com/2014/09/nasa-boeing-space-x-iss.html?
    utm_source=dlvr.it&utm_medium=twitter
```

Listing 9: uri_frequencies file