

CS532 Web Science: Assignment 4

Finished on February 25, 2016

Dr. Michael L. Nelson

Naina Sai Tipparti
ntippart@cs.odu.edu

Contents

Problem 1	2
Question	2
Answer	2
Problem 2	7
Question	7
Answer	7
Problem 4	11
Question	11
Answer	11

Listings

1	mln.graphml	2
2	get_graphml.py	4
3	Sort command	4
4	Graph Creation Script	4
5	get_followers.py	7
6	Sort command	9
7	Graph Creation Script for Twitter	9
8	get_followers.py	11
9	get_friends.py	11

List of Figures

1	Sample output of number of friends	3
2	Sample output of number of friends	3
3	The Friendship Graph for Facebook	5
4	The Friendship Graph for Twitter Followers	10
5	Sample output of number of following	12
6	The Friendship Graph for Twitter Following	13

List of Tables

1	Statistics on the count of Dr. Nelson Facebook Friends' Friends, values straight from R	6
2	Statistics on the count of Dr.Nelson Followers, values straight from R	10
3	Statistics on the count of Dr.Nelson Following, values straight from R	12

Problem 1

Question

Determine if the friendship paradox holds for my Facebook account. Compute the mean, standard deviation, and median of the number of friends that my friends have. Create a graph of the number of friends (y-axis) and the friends themselves, sorted by number of friends (x-axis). (The friends don't need to be labeled on the x-axis: just f1, f2, f3, ... fn.) Do include me in the graph and label me accordingly.

This used to be more interesting when you could more easily download your friend's friends data from Facebook. Facebook now requires each friend to approve this operation, effectively making it impossible.

I will email to the list the XML file that contains my Facebook friendship graph ca. Oct, 2013. The interesting part of the file looks like this (for 1 friend):

```
<node id="Johan_Bollen_1448621116">
  <data key="Label">Johan Bollen</data>
  <data key="uid"><![CDATA[1448621116]]></data>
  <data key="name"><![CDATA[Johan Bollen]]></data>
  <data key="mutual_friend_count"><![CDATA[37]]></data>
  <data key="friend_count"><![CDATA[420]]></data>
</node>
```

Listing 1: mln.graphml

It is in GraphML format: <http://graphml.graphdrawing.org/>

Answer

A Python program, *get_graphml.py*, has been written to extract number of friends that Dr. Nelson friends have. The program will search for this information in a file called *mln.graphml*. The output of this program would be like the following:

```

Naina Sai Tipparti@DESKTOP-2FU7AJC ~/a4/q1
$ python graphml.py
Friends-count Friend-screen-name
13 Simeon Warner
244 Simeon Warner
17 Drew Munro
575 Drew Munro
12 Mat Kelly
421 Mat Kelly
1 Benjamin Lok
539 Benjamin Lok
8 Camden Elliott Matherne
784 Camden Elliott Matherne
4 Barbara Burns Moran
317 Barbara Burns Moran
13 Jewel Ward
448 Jewel Ward
36 Geneva Henry
236 Geneva Henry
31 Timothy DiLauro
561 Timothy DiLauro
20 Maria Lugo
833 Maria Lugo
14 Frank McCown
752 Frank McCown
2 Hollie Chessman
763 Hollie Chessman
13 Sally Jo Cunningham
155 Sally Jo Cunningham
13 Leslie Carr
195 Leslie Carr
2 James Florance

```

Figure 1: Sample output of number of friends

```

77 Dale Andrews
8 Janet Schultz Brunelle
308 Janet Schultz Brunelle
13 Justin F. Brunelle
415 Justin F. Brunelle
2 Trey Arthur
111 Trey Arthur
18 Dave White
328 Dave White
2 Mike Koch
123 Mike Koch
6 Joel Carter
104 Joel Carter
9 Moustafa Aly
538 Moustafa Aly
7 Paul Ayris
147 Paul Ayris
1 Steve Bayer
353 Steve Bayer
3 Brooks Childers
59 Brooks Childers
2 Mary McManus
41 Mary McManus
1 Ashley Song
96 Ashley Song
3 Dongwon Lee
85 Dongwon Lee
4 Winnie Elliott
25 Winnie Elliott
8 Thomas Allen
39 Thomas Allen

```

Number of Dr. Nelson's friends ,who allow to retrieve their friends count, is 319 out of 165

Figure 2: Sample output of number of friends

```

1  # -*- encoding: utf-8 -*-
  #!/usr/bin/python
  from __future__ import unicode_literals
  import xml.etree.cElementTree as et
  from bs4 import BeautifulSoup
6  from urlparse import parse_qs
  import unicodedata
  import urllib2
  import re
  import os
11 import sys

  print '%-15s %-20s' %('Friends-count', 'Friend-screen-name')

  file = "mln.graphml"
16 handler = open(file).read()
  soup = BeautifulSoup(handler)
  i = 0
  all = 0
  for message in soup.find_all('node'):
21     all += 1
     foo = et.XML(str(message))
     name = ''
     for e in foo:
         if ('graphml_count' in str(e.items())):
26             print '%-15s %-20s' % (e.text, name)
             with open('friend_counts', 'a') as outfile:
                 outfile.write('%-15s %-20s\n' % (e.text, name))
             i += 1
         if ('name' in str(e.items())):
31             name = e.text
  print "\nNumber of Dr. Nelson's friends ,who allow to retrieve their friends count, is "+str
    (i)+" out of "+str(all)

```

Listing 2: get_graphml.py

I would like to let you know that even though Dr. Nelson have 319 friends, only 165 allow me to see their number of friends. This will affect the statistical result. For example, instead of dividing by 319 to get the mean, we divide by 165.

The `graphml_counts` file was ordered in place with the Unix command in Listing 3.

```
Naina Sai Tipparti@DESKTOP-2FU7AJC ~/a4/q1 cat graphml_counts | sort -g -o graphml_counts
```

Listing 3: Sort command

This file was then processed by the R script shown in Listing 4 to produce the graph in Figure 3

```

#!/usr/bin/Rscript

# read data
4 data <- read.table('D:/cs532/a4/q1/graphml_counts', sep=",")
x <- seq(1, length(data$V1))
y <- data$V1

# get notable values
9 mln_idx <- grep("phonedude_mln", data$V2)
med_val <- median(data$V1)
med_idx <- which(abs(y - med_val) == min(abs(y - med_val)))
mean_val <- mean(data$V1)
mean_idx <- which(abs(y - mean_val) == min(abs(y - mean_val)))

```

```

14 std_dev <- sd(data$V1)

# draw the graph
pdf("D:/cs532/a4/q1/facebook_graphml.pdf")
plot(x, y, type="l", log="y", pch=19, main="Dr. Nelson's Friends' Friends",
19 ylab="Number of Friends", xlab="Index of Friend")

# illustrate points of interest
abline(h=data$V1[mln_idx], col="red")

24 # The Legend of the Data
legend(x=82, y=5, cex=0.8, lty=c(1, 1),
      col=c("red", "white", "white", "white", "white"),
      c(paste("Nelson: ", data$V1[mln_idx]), paste("median: ", med_val),
        paste("mean: ", format(round(mean_val, 4), nsmall = 4)),
29       paste("std dev: ", format(round(std_dev, 4), nsmall = 4)),
        paste("median + 1 std dev: ", format(round(med_val + std_dev, 4), nsmall = 4))))
dev.off()

```

Listing 4: Graph Creation Script

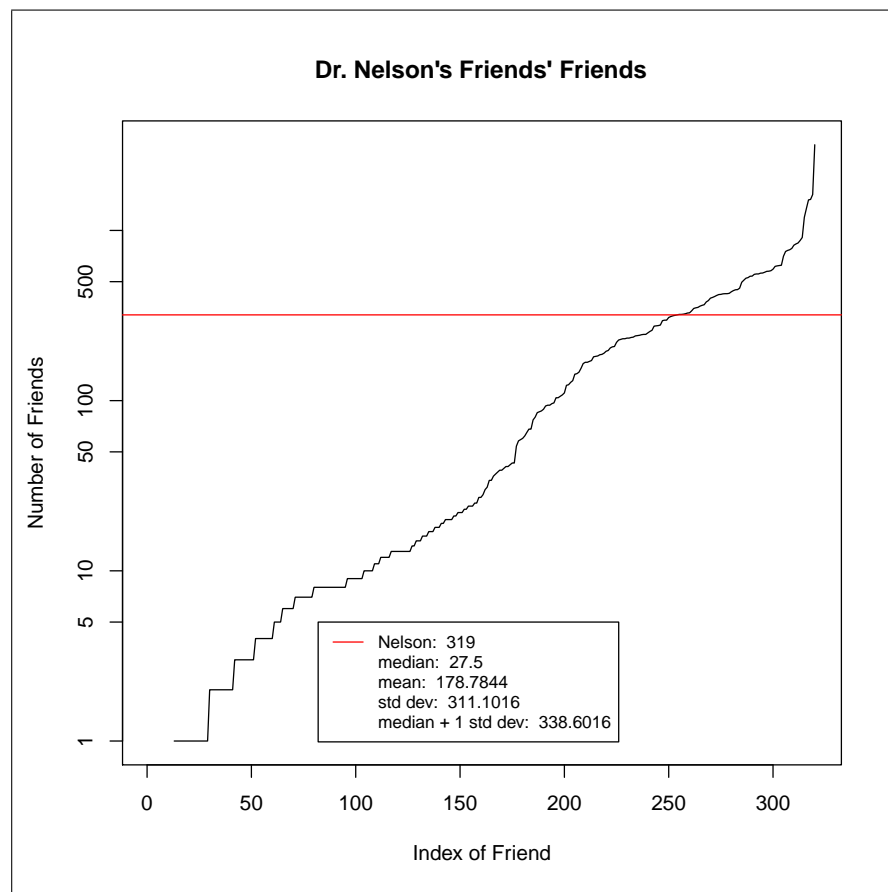


Figure 3: The Friendship Graph for Facebook

Mean	178.7844
Median	27.5
Std Dev	311.1016

Table 1: Statistics on the count of Dr. Nelson Facebook Friends' Friends, values straight from R

The median, mean and standard deviation were all calculated, with the median, mean and median plus one standard deviation.

Problem 2

Question

Determine if the friendship paradox holds for your Twitter account. Since Twitter is a directed graph, use “followers” as value you measure (i.e., “do your followers have more followers than you?”).

Generate the same graph as in question #1, and calculate the same mean, standard deviation, and median values.

For the Twitter 1.1 API to help gather this data, see:

<https://dev.twitter.com/docs/api/1.1/get/followers/list>

If you do not have followers on Twitter (or don't have more than 50), then use my twitter account “phonedude_mln”.

Answer

Using Dr. Michael Nelson's Twitter account and the Twitter API [?], specifically the GET friends/list [?] request, all of Dr. Nelson's Twitter friends were obtained and saved to the file called `followers`. This method also uses the API's paginating scheme: when there are a large number of results for a query, the API will send a cursor index to show that there are more results to process and that more requests are needed. The code to do this is in Listing 5.

```
# -*- encoding: utf-8 -*-
from __future__ import unicode_literals
import re
4 import os
import sys
import json
import requests
import subprocess
9 from urlparse import parse_qs
from requests_oauthlib import OAuth1

REQUEST_TOKEN_URL = "https://api.twitter.com/oauth/request_token"
AUTHORIZE_URL = "https://api.twitter.com/oauth/authorize?oauth_token="
14 ACCESS_TOKEN_URL = "https://api.twitter.com/oauth/access_token"

CONSUMER_KEY = "IA3ACTYCPDE8G5rNYFBMI1hMm"
CONSUMER_SECRET = "4zGZNDRA2m32dsq7nCMfwJojGSanz6ohgf4ZaNWKDxCaabPUai"
OAUTH_TOKEN = "798668178-bH8DbMpNuWkfhAHxuODgWSHwQE65B1WZnc4Ahtej"
19 OAUTH_TOKEN_SECRET = "FhykPKnQcgKQBE43os2bDZ31ugH9RVSG3HYoOL7QG7RNC"

def setup_oauth():
    """Authorize your app via identifier."""
    # Request token
24    oauth = OAuth1(CONSUMER_KEY, client_secret=CONSUMER_SECRET)
    r = requests.post(url=REQUEST_TOKEN_URL, auth=oauth)
    credentials = parse_qs(r.content)

    resource_owner_key = credentials.get('oauth_token')[0]
29    resource_owner_secret = credentials.get('oauth_token_secret')[0]

    # Authorize
```



```

    authorize_url = AUTHORIZE_URL + resource_owner_key
    print 'Please go here and authorize: ' + authorize_url

34
    verifier = raw_input('Please input the verifier: ')
    oauth = OAuth1(CONSUMER_KEY,
                   client_secret=CONSUMER_SECRET,
                   resource_owner_key=resource_owner_key,
39                   resource_owner_secret=resource_owner_secret,
                   verifier=verifier)

    # Finally, Obtain the Access Token
    r = requests.post(url=ACCESS_TOKEN_URL, auth=oauth)
44    credentials = parse_qs(r.content)
    token = credentials.get('oauth_token')[0]
    secret = credentials.get('oauth_token_secret')[0]

    return token, secret

49

def get_oauth():
    oauth = OAuth1(CONSUMER_KEY,
                   client_secret=CONSUMER_SECRET,
54                   resource_owner_key=OAUTH_TOKEN,
                   resource_owner_secret=OAUTH_TOKEN_SECRET)

    return oauth

if __name__ == "__main__":
59     if not OAUTH_TOKEN:
        token, secret = setup_oauth()
        print "OAUTH_TOKEN: " + token
        print "OAUTH_TOKEN_SECRET: " + secret
        print
64     else:
        twitterUser = "phonedude_mln"

        print 'Searching Twitter for followers counts of '+twitterUser+'s followers: '
        oauth = get_oauth()

69
        print '%-15s %-20s' %('Followers_count', 'Follower-screen-name')

        # initial reading from the twitter account where cursor = -1 (e.g. first page)
        r = requests.get(url="https://api.twitter.com/1.1/followers/list.json?cursor=-1&
                           count=2000&screen_name="+twitterUser+"&skip_status=true&include_user_entities=
                           false", auth=oauth)
74        counter = 0
        res = r.json()
        while True:
            raw_res = res['users']
            for init_url in raw_res:
79                counter = counter + 1
                print '%-15d %-20s' %(init_url['followers_count'], init_url['screen_name']).
                    encode('ascii', 'replace'))
                with open('friend_counts', 'a') as outfile:
                    outfile.write('%-15d %-20s\n' %(init_url['followers_count'], init_url['
                        screen_name']).encode('ascii', 'replace'))
            if str(res['next_cursor']) == '0':
84                break
            else:
                r = requests.get(url="https://api.twitter.com/1.1/followers/list.json?cursor
                                   ="+str(res['next_cursor'])+"&count=100&screen_name="+twitterUser+"&
                                   skip_status=true&include_user_entities=false", auth=oauth)
                res = r.json()

89 print '\nNumber of '+twitterUser+'s followers is: "+str(counter)

```

Listing 5: get_followers.py

To reduce the impact of high HTTP traffic, the Twitter API rate-limits most requests – the one needed to obtain a user’s friends list has a limit of fifteen message per fifteen minutes. Any requests received from a user or service that has reached the limit will be denied. The friends of Dr. Nelson’s friends were then obtained with the same `get_followers` method from Listing 5 and stored in a file called `followers_counts`, each on a single line preceded by their friend count. All of these operations were controlled by a main method, which is shown in Listing 5.

The `followers_counts` file was ordered in place with the Unix command in Listing 6.

```
1 Naina Sai Tippiarti@DESKTOP-2FU7AJC ~/a4/q3 cat followers_counts | sort -g -o
  followers_counts
```

Listing 6: Sort command

This file was then processed by the R script shown in Listing 7 to produce the graph in Figure 4

```
#!/usr/bin/Rscript

# read data
4 data <- read.table('D:/cs532/a4/q2/followers_counts')
  x <- seq(1, length(data$V1))
  y <- data$V1

# get notable values
9 mln_idx <- grep("phonedude_mln", data$V2)
  med_val <- median(data$V1)
  med_idx <- which(abs(y - med_val) == min(abs(y - med_val)))
  mean_val <- mean(data$V1)
  mean_idx <- which(abs(y - mean_val) == min(abs(y - mean_val)))
14 std_dev <- sd(data$V1)

# draw the graph
pdf("D:/cs532/a4/q2/followers_plot.pdf")
plot(x, y, type="l", log="y", pch=19, main="Dr. Nelson's Friends' Friends",
19 ylab="Number of Friends", xlab="Index of Friend")

# illustrate points of interest
abline(h=data$V1[mln_idx], col="red")

24 # The Legend of the Data
  legend(x=82, y=5, cex=0.8, lty=c(1, 1),
        col=c("red", "white", "white", "white", "white"),
        c(paste("Nelson: ", data$V1[mln_idx]), paste("median: ", med_val),
          paste("mean: ", format(round(mean_val, 4), nsmall = 4)),
          paste("std dev: ", format(round(std_dev, 4), nsmall = 4)),
          paste("median + 1 std dev: ", format(round(med_val + std_dev, 4), nsmall = 4))))
29 dev.off()
```

Listing 7: Graph Creation Script for Twitter

Mean	1024.7424
Median	258
Std Dev	4135.7974

Table 2: Statistics on the count of Dr.Nelson Followers, values straight from R

The median, mean and standard deviation were all calculated, with the median, mean and median plus one standard deviation.




Figure 4: The Friendship Graph for Twitter Followers

Problem 4

Question

Repeat question #2, but change “follower” to “following”? In other words, are the people I am following following more people?

Answer

Same instructions explained in question 2 are used to answer this question. As I remember, only one change has been made to *get_followers.py*.

- The old request:

```
r = requests.get(url="https://api.twitter.com/1.1/followers/list.json?cursor=-1&count=2000&screen_name="+twitterUser+"&skip_status=true&include_user_entities=false", auth=oauth)
```

Listing 8: *get_followers.py*

- After modifying:

```
r = requests.get(url="https://api.twitter.com/1.1/friends/list.json?cursor=-1&count=2000&screen_name="+twitterUser+"&skip_status=true&include_user_entities=false", auth=oauth)
```

Listing 9: *get_friends.py*

All new changes are stored in a new file called *get_friends.py*. The following is the output after running the Python program:

```
Naina Sai Tipparti@DESKTOP-2F07A3C ~/a4/q4
$ python get_friends.py
Searching Twitter for following counts of phonedude_m1n's following:
Friends_count  Friends-screen-name
873            BergisJules
1              stormcrawlerapi
1              Sci_Hub
71             stewartbrand
133            documentnow
4              DanMilanko
30             baconmeteor
1229           nowiskie
74            FarbThink
243            whanamura
793            coachwiles
250            TKP_Recruiting
323            anarchist
2296           archivesnext
260            CoachFuente
325            NtelsBr
1              IPres2016
855            YourPaperSucks
533            e_mccain
512            dmimno
1              UberFacts
351            thekeyplay
38            vneblitt
13            acnwala
54            maturbani
28            kaylamar1e0110
57            Jcdl2016
144            Clarellwellyn
433            substack
```

Figure 5: Sample output of number of following

Mean	484.9214
Median	225
Std Dev	729.5275

Table 3: Statistics on the count of Dr.Nelson Following, values straight from R




Figure 6: The Friendship Graph for Twitter Following

I think here also the friendship paradox holds for Dr. Nelson Twitter account (for following) since the mean is equal to 484.9214 which is greater than most of the number of following of Dr. Nelson following including himself.

