

# CS532 Web Science: Assignment 7

Finished on March 31, 2016

*Dr. Michael L. Nelson*

**Naina Sai Tipparti**  
ntippart@cs.odu.edu

## Contents

<b>Problem 1</b>	<b>3</b>
Question . . . . .	3
1.1 Answer . . . . .	3
1.1.1 3-Users Closest to Me in Terms of Age, Gender, and Occupation . . . .	5
1.1.2 3-Users Closest to Me Top 3 Favorite Films . . . . .	5
1.1.3 3-Users Closest to Me Least 3 Favorite Films . . . . .	6
1.1.4 Substitute Me . . . . .	7
<b>Problem 2</b>	<b>8</b>
Question . . . . .	8
2.1 Answer . . . . .	8
2.1.1 5-Users Most and Least Correlated to the Substitute Me . . . . .	8
<b>Problem 3</b>	<b>10</b>
Question . . . . .	10
Answer . . . . .	10
3.1.1 Top 5 Recommendations of Unseen films for Substitute Me . . . . .	11
3.1.2 Bottom 5 Recommendations of Unseen films for Substitute Me . . . . .	11
<b>Problem 4</b>	<b>12</b>
Question . . . . .	12
Answer . . . . .	12
<b>Problem 5</b>	<b>14</b>
Question . . . . .	14
Answer . . . . .	14
<b>Problem 6</b>	<b>15</b>
Question . . . . .	15
Answer . . . . .	15
<b>Appendix A</b>	<b>16</b>

## Listings

1	tabulate function . . . . .	4
2	Question 1 code . . . . .	4
3	get_top functions . . . . .	4
4	flatten function . . . . .	4
5	Question 2 code . . . . .	8
6	sim_pearson function . . . . .	8
7	Question 3 code . . . . .	10
8	movies_not Rated function . . . . .	10

9	get_avg functions . . . . .	12
10	Question 4 code . . . . .	12
11	pearson movie function . . . . .	12
12	all code used . . . . .	16

## List of Tables

1	All-Users Closest Match . . . . .	5
2	Top 3 Favorite Films . . . . .	5
3	Least 3 Favorite Films . . . . .	6
4	User <a href="#">135</a> was selected as substitute me . . . . .	7
5	Most Correlated Substitute Me . . . . .	9
6	Least Correlated Substitute Me . . . . .	9
7	Top 5 unseen movies recommendations . . . . .	11
8	Least 5 unseen movies recommendations . . . . .	11
9	Most Favourite Movie <a href="#">Godfather, The (1972)</a> Correlated . . . . .	13
10	Least Favourite Movie <a href="#">Godfather, The (1972)</a> Correlated . . . . .	13

# Problem 1

## Question

The goal of this project is to use the basic recommendation principles we have learned for user-collected data. You will modify the code given to you which performs movie recommendations from the MovieLens data sets.

The MovieLens data sets were collected by the GroupLens Research Project at the University of Minnesota during the seven-month period from September 19th, 1997 through April 22nd, 1998. We are using the “100k dataset”; available for download from:

<http://grouplens.org/datasets/movielens/100k/>

The code for reading from the `u.data` and `u.item` files and creating recommendations is described in the book *Programming Collective Intelligence*. Feel free to modify the PCI code to answer the following questions.

Find 3 users who are closest to you in terms of age, gender, and occupation. For each of those 3 users:

- what are their top 3 favorite films?
- bottom 3 least favorite films?

Based on the movie values in those 6 tables (3 users X (favorite + least)), choose a user that you feel is most like you. Feel free to note any outliers (e.g., “I mostly identify with user 123, except I did not like ‘Ghost’ at ll”).

This user is the “substitute you”.

### 1.1 Answer

Each question was answered by using some combination of the existing functions from `recommendations.py` and some functions that were added. All of the tables provided were created using the `tabulate` function (with some minor edits), which is shown in Listing 1. All of the code used can be found in Appendix A, which contains Listing 12.

Question 1 was solved using the code in Listing 2, which utilizes the added `flatten` and `get_top` functions, which are found in Listing 3 and 4.

The results for Question 1 are shown in Tables 1, 2, 3 and 4.

```

def tabulate(tuples, caption, label, colnames, output):
    output.write('\begin{table}[h!]\n')
    output.write('\centering\n')
    opts = '|' + '|' * len(tuples[0]) + '|'
    290 output.write('\begin{tabular}{{{0}}}\n'.format(opts))
    output.write('\hline\n')
    header = '&'.join(['{}' for i in xrange(len(tuples[0]))]).format(*colnames)
    output.write(header + '\\\\n\\hline\n')
    295 for item in tuples:
        temp = '&'.join(['{}' for i in xrange(len(item))])
        output.write(temp.format(*item) + '\\\\n\n')
    output.write('\hline\n\end{tabular}\n')
    output.write('\caption{{{0}}}\n'.format(caption))
    output.write('\label{tab:{0}}\n'.format(label))
    300 output.write('\end{table}\n\n')

```

Listing 1: tabulate function

```

getuser = {}
user_filter = lambda x: x['gender'] == 'M' and x['job'] == 'student' and int(x['age']
    ]) == 23
ratings = []
385 for mid, movie in movies.iteritems():
    for user, user_ratings in prefs.iteritems():
        if user_filter(users[user]) and user_ratings.has_key(movies[mid]):
            getuser.setdefault(int(user), {})
            getuser[int(user)][movies[mid]] = float(user_ratings[movies[mid]])
390 ratings.append(user_ratings[movies[mid]])
sorted_getuser = {}
movie_sort = {}
for user, user_movie in getuser.items():
    user_movie_sort = sorted(user_movie.items(), key=itemgetter(1), reverse=False)
395 for title, rating in user_movie_sort[:3]:
    movie_sort.setdefault(title, rating)
    # print movie_sort
sorted_getuser.setdefault(user, user_movie_sort[:3])
sorted_avg_all = sorted(sorted_getuser.items(), key=itemgetter(0), reverse=False)
400 # print sorted_avg_all
top_raters = get_top(sorted_avg_all, key=lambda x, i: x[i][1][0])
top_raters = [flatten(rater) for rater in sorted_avg_all]
tabulate(top_raters, 'Users', 'user', ('User', 'Rating', 'Movie'), outfile)
print "done with 1"

```

Listing 2: Question 1 code

```

def get_top(sorted_list, key=lambda x, i: x[i][1], n=5):
    top = key(sorted_list, 0)
    top_items = []
    265 i = 0
    while i < n or key(sorted_list, i) == top:
        top_items.append(sorted_list[i])
        if i < n and key(sorted_list, i) != top:
            top = key(sorted_list, i)
        i += 1
    270 return top_items

```

Listing 3: get\_top functions

```

def flatten(tup, f=lambda x: (x[0], x[1][0][1], x[1][0][0])):
    return f(tup)

```

Listing 4: flatten function

**1.1.1 3-Users Closest to Me in Terms of Age, Gender, and Occupation**

Id	Age	Gender	Occupation
33	23	M	student
37	23	M	student
66	23	M	student
135	23	M	student
391	23	M	student
408	23	M	student
706	23	M	student
838	23	M	student

Table 1: All-Users Closest Match

**1.1.2 3-Users Closest to Me Top 3 Favorite Films**

User	Rating	Movie
33	5.0	Titanic (1997)
	4.0	Game, The (1997)
	4.0	Air Force One (1997)
37	5.0	Pulp Fiction (1994)
	5.0	Raiders of the Lost Ark (1981)
	5.0	Terminator, The (1984)
66	5.0	Return of the Jedi (1983)
	5.0	Air Force One (1997)
	5.0	Ransom (1996)
135	5.0	Silence of the Lambs, The (1991)
	4.0	Rear Window (1954)
	4.0	Liar Liar (1997)
391	5.0	Rear Window (1954)
	5.0	Magnificent Seven, The (1954)
	5.0	Blues Brothers, The (1980)
408	5.0	Liar Liar (1997)
	5.0	Lost Highway (1997)
	5.0	Everyone Says I Love You (1996)
706	5.0	Phenomenon (1996)
	5.0	Edge, The (1997)
	5.0	Star Wars (1977)
838	5.0	Bringing Up Baby (1938)
	5.0	Toy Story (1995)
	5.0	City of Lost Children, The (1995)

Table 2: Top 3 Favorite Films

**1.1.3 3-Users Closest to Me Least 3 Favorite Films**

User	Rating	Movie
33	3.0	Liar Liar (1997)
	3.0	Devil's Advocate, The (1997)
	3.0	Soul Food (1997)
37	1.0	Jurassic Park (1993)
	2.0	Twister (1996)
	2.0	Arrival, The (1996)
66	1.0	English Patient, The (1996)
	1.0	Muppet Treasure Island (1996)
	1.0	Excess Baggage (1997)
135	1.0	Tales from the Hood (1995)
	2.0	Jaws 2 (1978)
	2.0	Star Trek III: The Search for Spock (1984)
391	1.0	Mimic (1997)
	2.0	Star Trek: The Wrath of Khan (1982)
	2.0	Courage Under Fire (1996)
408	1.0	Mouse Hunt (1997)
	2.0	U Turn (1997)
	2.0	Conspiracy Theory (1997)
706	1.0	Game, The (1997)
	1.0	Fargo (1996)
	1.0	Crash (1996)
838	2.0	Mars Attacks! (1996)
	2.0	Independence Day (ID4) (1996)
	2.0	Air Force One (1997)

Table 3: Least 3 Favorite Films

### 1.1.4 Substitute Me

Selecting substitute me was a difficult choice. I couldn't relate to user 33, 37, 66, 391, 706, 838, despite the fact that I would rank "Titanic (1997)" similarly with a 5.

I am greatly removed from user 33, 37, 66, 391, 706, 838 all the most maximum ranked movies by 33, 37, 66, 391, 706, 838 would be in my minimum ranked movies determination.

I am not that nearby either with user 135, 408, the most elevated ranked movies will be in my 4 rate selection. So, User 135 was selected as substitute of me.

User	Rating	Movie
135	5.0	Silence of the Lambs, The (1991)
	4.0	Rear Window (1954)
	4.0	Liar Liar (1997)
	1.0	Tales from the Hood (1995)
	2.0	Jaws 2 (1978)
	2.0	Star Trek III: The Search for Spock (1984)

Table 4: User 135 was selected as substitute me



## Problem 2

### Question

Which 5 users are most correlated to the substitute you? Which 5 users are least correlated (i.e., negative correlation)?

### 2.1 Answer

#### 2.1.1 5-Users Most and Least Correlated to the Substitute Me

Question 2 was solved using the code shown in Listing 5, with the `sim_pearson` and the `get_top` functions from Listings 6 and 3, respectively. The `flatten` function from Listing 4 was used to flattened the oddly arranged tuple that was created from the `sim_pearson` function.

```

405     most_correlated = {}
        least_correlated = {}
        for user, rest in users.iteritems():
            if int(user) == 135:
                for user1, rest in users.iteritems():
410                     if user == user1:
                        pass
                    else:
                        r = sim_pearson(prefs, user, user1)
                        if r == 1.0:
415                             most_correlated.setdefault(int(user1), r)
                        if r == -1.0:
                            least_correlated.setdefault(int(user1), r)
        sorted_most_correlated = sorted(most_correlated.items(), key=itemgetter(0), reverse=
            False)[len(most_correlated) - 6:-1]
        sorted_least_correlated = sorted(least_correlated.items(), key=itemgetter(0),
            reverse=False)[len(least_correlated) - 6:-1]
420     tabulate(sorted_most_correlated, 'Most Correlated', 'most', ('User', 'Pearson\'s r'),
        , outfile)
        tabulate(sorted_least_correlated, 'Least Correlated', 'least', ('User', 'Pearson\'s
            r'), outfile)
        print "done with 2"
```

Listing 5: Question 2 code

```

45 def sim_pearson(prefs, p1, p2):
    # Get the list of mutually rated items
    si={}
    for item in prefs[p1]:
        if item in prefs[p2]:
50         si[item]=1

    # if they are no ratings in common, return 0
    if len(si)==0: return 0

55     # Sum calculations
    n=len(si)

    # Sums of all the preferences
    sum1=sum([prefs[p1][it] for it in si])
    sum2=sum([prefs[p2][it] for it in si])
60

    # Sums of the squares
```

```

sum1Sq=sum([pow(prefs[p1][it],2) for it in si])
sum2Sq=sum([pow(prefs[p2][it],2) for it in si])

# Sum of the products
pSum=sum([prefs[p1][it]*prefs[p2][it] for it in si])

# Calculate r (Pearson score)
num=pSum-(sum1*sum2/n)
den=sqrt((sum1Sq-pow(sum1,2)/n)*(sum2Sq-pow(sum2,2)/n))
if den==0: return 0

r=num/den
filename = p1
outfile = open(filename, 'a')
outfile.write(str(r) + "\t\t" + str(user1))
outfile.write("\n")
outfile.close()
return r

```

Listing 6: sim\_pearson function

The results for Question 2 are shown in Table 5 and 6.

User	Pearson's r
20	1.0
74	1.0
161	1.0
219	1.0
284	1.0

Table 5: Most Correlated Substitute Me

User	Pearson's r
26	-1.0
39	-1.0
55	-1.0
241	-1.0
300	-1.0

Table 6: Least Correlated Substitute Me

## Problem 3

### Question

Compute ratings for all the films that the substitute you hasn't seen. Provide a list of the top 5 recommendations for films that the substitute you should see. Provide a list of the bottom 5 recommendations (i.e., films the substitute you is almost certain to hate).

### Answer

Question 3 was solved using the code shown in Listing 7, with the `movies_notRated` function from Listings 8, respectively. The outcome is put away in variable `movie_rated1` which will contain the ranking of movies not yet seen by `substitute me` sorted from slightest no doubt sought recommended preference to the probably recommended preference.

The results for Question 3 are shown in Table 7 and 8.

```

425 movie_rated1 = {}
    for user, rest in users.iteritems():
        if int(user) == 135:
            movie_rated1 = movies_notRated(prefs, movies, user, genre)
            sorted_most_recomm = sorted(movie_rated1.items(), key=itemgetter(1), reverse=False)[
                len(movie_rated1) - 6:-1]
            sorted_least_recomm = sorted(movie_rated1.items(), key=itemgetter(1), reverse=True)[
                len(movie_rated1) - 6:-1]
            tabulate(sorted_most_recomm, 'Top 5 unseen movies recommendations', 'recommtop', (
                'Title', 'Recomm Ranking'), outfile)
430 tabulate(sorted_least_recomm, 'Least 5 unseen movies recommendations', 'recommleast',
            ('Title', 'Recomm Ranking'), outfile)
    print "done with 3"
```

Listing 7: Question 3 code

```

def movies_notRated(prefs, movies, userid, movie_type):
    userRatings = prefs[userid]
    count = 0
    movie_rated = {}
305 for id, title in movies.items():
        temp = 0
        for item, rating in userRatings.items():
            if title == item:
310 temp = 1
                movie_rated.setdefault(title, rating)
        for name, genre in movie_type.items():
            if title == name:
                if temp == 0:
315 try:
                    if genre == 'comedy':
                        movie_rated[title] = '3.0'
                        temp = 0
                    elif genre == 'action':
320 movie_rated[title] = '3.0'
                        temp = 0
                    elif genre == 'crime':
                        movie_rated[title] = '3.0'
                        temp = 0
325 elif genre == 'drama':
                        movie_rated[title] = '4.0'
                        temp = 0
```

```

330         elif genre == 'thriller':
331             movie Rated[t title] = '5.0'
332             temp = 0
333         elif genre == 'scifi':
334             movie Rated[t title] = '4.0'
335             temp = 0
336         elif genre == 'war':
337             movie Rated[t title] = '3.0'
338             temp = 0
339         elif genre == 'mystery':
340             movie Rated[t title] = '3.0'
341             temp = 0
342         else:
343             movie Rated[t title] = '1.0'
344             temp = 0
345     except:
346         pass
347     return movie Rated

```

Listing 8: movies\_not Rated function

### 3.1.1 Top 5 Recommendations of Unseen films for Substitute Me

Title	Ranking
Stranger in the House (1997)	5.0
Killer (Bulletproof Heart) (1994)	5.0
Gaslight (1944)	5.0
Office Killer (1997)	5.0
Unforgettable (1996)	5.0

Table 7: Top 5 unseen movies recommendations

### 3.1.2 Bottom 5 Recommendations of Unseen films for Substitute Me

Title	Ranking
Jaws 2 (1978)	2.0
Star Trek III: The Search for Spock (1984)	2.0
Highlander (1986)	2.0
Hard Target (1993)	2.0
Tales From the Crypt Presents: Demon Knight (1995)	2.0

Table 8: Least 5 unseen movies recommendations

## Problem 4

### Question

Choose your (the real you, not the substitute you) favorite and least favorite film from the data. For each film, generate a list of the top 5 most correlated and bottom 5 least correlated films. Based on your knowledge of the resulting films, do you agree with the results? In other words, do you personally like / dislike the resulting films?

### Answer

Question 4 was solved using the code shown in Listing 10, with the `pearson` and the `get_avg` functions from Listings 11 and 9, respectively. The `get_avg` function uses the `mean` function from the `numpy` python library [1].

The results for Question 4 are shown in Table 9 and 10.

```

def get_avg(prefs, mid, user_filter=lambda x: True):
    ratings = []
    for user, user_ratings in prefs.iteritems():
        if user_filter(users[user]) and user_ratings.has_key(movies[mid]):
            ratings.append(user_ratings[movies[mid]])
    if not ratings:
        return 0.0
    return mean(ratings)

```

Listing 9: `get_avg` functions

```

fav_cor = {}
nfav_cor = {}
avg1 = get_avg(prefs, mfavs)
for id, title in movies.items():
    avg2 = get_avg(prefs, id)
    r = pearson(avg1, avg2, prefs, mfavs, title)
    if r == 1.0:
        fav_cor.setdefault(str(title), str(r))
    elif r == -1.0:
        nfav_cor.setdefault(str(title), str(r))
sorted_fav_correlated = sorted(fav_cor.items(), key=itemgetter(0), reverse=False)[
    len(fav_cor) - 6:-1]
sorted_nfav_correlated = sorted(nfav_cor.items(), key=itemgetter(0), reverse=False)[
    len(nfav_cor) - 6:-1]
tabulate(sorted_fav_correlated, 'Most Favourite Movie \textcolor{blue}{Godfather,
    The (1972)} Correlated', 'fav', ('Title', 'Pearson\'s r'), outfile)
tabulate(sorted_nfav_correlated, 'Least Favourite Movie \textcolor{blue}{Godfather,
    The (1972)} Correlated', 'nfav', ('Title', 'Pearson\'s r'), outfile)
print "done with 4"

```

Listing 10: Question 4 code

```

def pearson(avg1, avg2, prefs, movie1, movie2, movies):
    r = {}
    r = 0
    top = 0

```

```

    rbtm = 0
    lbtm = 0

    for user, movie in prefs.items():
        userRatings = prefs[user]
        for item, rating in userRatings.items():
            if movie1 == movies[item]:
                for user1, move in prefs.items():
                    userRatings = prefs[user1]
                    for item1, rating1 in userRatings.items():
                        if movie2 == movies[item1]:
                            if user1 == user:
                                top += (rating-avg1)*(rating1-avg2)
                                lbtm += (rating-avg1)*(rating-avg1)
                                rbtm += (rating1-avg2)*(rating1-avg2)

    if lbtm == 0 or rbtm == 0:
        r = 5
        return r
    else:
        r = top/(math.sqrt(lbtm)* math.sqrt(rbtm))
    return r

```

Listing 11: pearson movie function

Title	Person's r
Gang Related (1997)	1.0
Love and Other Catastrophes (1996)	1.0
Rhyme & Reason (1997)	1.0
When We Were Kings (1996)	1.0
Dark City (1998)	1.0

Table 9: Most Favourite Movie [Godfather, The \(1972\)](#) Correlated

Title	Person's r
Paris, France (1993)	-1.0
Hugo Pool (1997)	-1.0
Naked in New York (1994)	-1.0
Turbo: A Power Rangers Movie (1997)	-1.0
Heavyweights (1994)	-1.0

Table 10: Least Favourite Movie [Godfather, The \(1972\)](#) Correlated

The recommendations results are not amazing to me. As expressed in User [135](#) was selected as substitute me4 arrangement user [135](#) is nearest to me of course. Be that as it may, the outcomes are close in wording how I feel about **substitute me** ranking. I hate all movies not recommended and the recommended ones I would rate them somewhere around 2 and 3, like the most astounding ranking for user [135](#).

## Problem 5

### Question

Rank the 1,682 movies according to the 1997/1998 MovieLense data. Now rank the same 1,682 movies according to today's (March 2016) IMDB data (break ties based on # of users, for example: 7.2 with 10,000 raters  $>$  7.2 with 9,000 raters).

Draw a graph, where each dot is a film (i.e., 1,682 dots). The x-axis is the MovieLense ranking and the y-axis is today's IMDB ranking.

What is Pearson's  $r$  for the two lists (along w/ the p-value)? Assuming the two user bases are interchangeable (which might not be a good assumption), what does this say about the attitudes about the films after nearly 20 years?

### Answer

## Problem 6

### Question

Repeat #6, but IMDB data from approximately July 31, 2005. What is the cumulative error (in days) from the desired target day of July 31, 2005? For example, if 1 memento is from July 1, 2005 and another memento is from July 31, 2006, then the cumulative error for the two mementos is 30 days + 365 days = 385 days.

Note: the URIs in the MovieLens data redirect, be sure to use the final values as URIs for the archives:

### Answer



## Appendix A

```

# A dictionary of movie critics and their ratings of a small
# set of movies
3 critics={'Lisa Rose': {'Lady in the Water': 2.5, 'Snakes on a Plane': 3.5,
    'Just My Luck': 3.0, 'Superman Returns': 3.5, 'You, Me and Dupree': 2.5,
    'The Night Listener': 3.0},
    'Gene Seymour': {'Lady in the Water': 3.0, 'Snakes on a Plane': 3.5,
    'Just My Luck': 1.5, 'Superman Returns': 5.0, 'The Night Listener': 3.0,
8    'You, Me and Dupree': 3.5},
    'Michael Phillips': {'Lady in the Water': 2.5, 'Snakes on a Plane': 3.0,
    'Superman Returns': 3.5, 'The Night Listener': 4.0},
    'Claudia Puig': {'Snakes on a Plane': 3.5, 'Just My Luck': 3.0,
    'The Night Listener': 4.5, 'Superman Returns': 4.0,
13    'You, Me and Dupree': 2.5},
    'Mick LaSalle': {'Lady in the Water': 3.0, 'Snakes on a Plane': 4.0,
    'Just My Luck': 2.0, 'Superman Returns': 3.0, 'The Night Listener': 3.0,
    'You, Me and Dupree': 2.0},
    'Jack Matthews': {'Lady in the Water': 3.0, 'Snakes on a Plane': 4.0,
18    'The Night Listener': 3.0, 'Superman Returns': 5.0, 'You, Me and Dupree': 3.5},
    'Toby': {'Snakes on a Plane': 4.5, 'You, Me and Dupree': 1.0, 'Superman Returns': 4.0}}

import math
from math import sqrt
23 from numpy import mean
from operator import itemgetter
from pprint import pprint
from sys import stdout

28 # Returns a distance-based similarity score for person1 and person2
def sim_distance(prefs, person1, person2):
    # Get the list of shared_items
    si={}
    for item in prefs[person1]:
33         if item in prefs[person2]: si[item]=1

    # if they have no ratings in common, return 0
    if len(si)==0: return 0

38    # Add up the squares of all the differences
    sum_of_squares=sum([pow(prefs[person1][item]-prefs[person2][item],2)
        for item in prefs[person1] if item in prefs[person2]])

    return 1/(1+sum_of_squares)
43

# Returns the Pearson correlation coefficient for p1 and p2
def sim_pearson(prefs, p1, p2):
    # Get the list of mutually rated items
    si={}
    for item in prefs[p1]:
48         if item in prefs[p2]:
            si[item]=1

    # if they are no ratings in common, return 0
53    if len(si)==0: return 0

    # Sum calculations
    n=len(si)

58    # Sums of all the preferences
    sum1=sum([prefs[p1][it] for it in si])
    sum2=sum([prefs[p2][it] for it in si])

    # Sums of the squares
63    sum1Sq=sum([pow(prefs[p1][it],2) for it in si])
    sum2Sq=sum([pow(prefs[p2][it],2) for it in si])

```

```

# Sum of the products
pSum=sum([ prefs[p1][it]*prefs[p2][it] for it in si])

68
# Calculate r (Pearson score)
num=pSum-(sum1*sum2/n)
den=sqrt((sum1Sq-pow(sum1,2)/n)*(sum2Sq-pow(sum2,2)/n))
if den==0: return 0

73
r=num/den
filename = p1
outfile = open(filename, 'a')
outfile.write(str(r) + "\t\t" + str(user1))
78
outfile.write("\n")
outfile.close()
return r

# Returns the best matches for person from the prefs dictionary.
83 # Number of results and similarity function are optional params.
def topMatches(prefs, person, n=5, similarity=sim_pearson):
    scores=[(similarity(prefs, person, other), other)
             for other in prefs if other!=person]
    scores.sort()
88
    scores.reverse()
    return scores[0:n]

# Gets recommendations for a person by using a weighted average
# of every other user's rankings
93 def getRecommendations(prefs, person, similarity=sim_pearson):
    totals={}
    simSums={}
    for other in prefs:
        # don't compare me to myself
        98 if other==person: continue
        sim=similarity(prefs, person, other)

        # ignore scores of zero or lower
        if sim<=0: continue
        103 for item in prefs[other]:

            # only score movies I haven't seen yet
            if item not in prefs[person] or prefs[person][item]==0:
                # Similarity * Score
                totals.setdefault(item,0)
                totals[item]+=prefs[other][item]*sim
                # Sum of similarities
                simSums.setdefault(item,0)
                simSums[item]+=sim

113
# Create the normalized list
rankings=[(total/simSums[item], item) for item, total in totals.items()]

# Return the sorted list
118 rankings.sort()
rankings.reverse()
return rankings

def transformPrefs(prefs):
123 result={}
for person in prefs:
    for item in prefs[person]:
        result.setdefault(item, {})

128
    # Flip item and person
    result[item][person]=prefs[person][item]
return result

def calculateSimilarItems(prefs, n=10, similarity=sim_distance):
133 # Create a dictionary of items showing which other items they

```

```

# are most similar to.
result={}
# Invert the preference matrix to be item-centric
itemPrefs=transformPrefs( prefs )
138 c=0
for item in itemPrefs:
    # Status updates for large datasets
    c+=1
    if c%100==0: print "%d / %d" % (c, len(itemPrefs))
143 # Find the most similar items to this one
    scores=topMatches( itemPrefs , item , n=n, similarity=similarity )
    result [ item]=scores
return result

148 def calcSimilarUsers( prefs , n=10, similarity=sim_distance ):
    result = {}
    itemPrefs = prefs
    c=0
    for item in itemPrefs:
153 c+=1
        if c%100==0: print "%d / %d" % (c, len(itemPrefs))
        scores = topMatches( itemPrefs , item , n=n, similarity=similarity )
        result [ item]=scores
    return result

158 def getRecommendedItems( prefs , itemMatch , user ):
    userRatings=prefs [ user ]
    scores={}
    totalSim={}
163 # Loop over items rated by this user
    for ( item , rating ) in userRatings.items( ):

        # Loop over items similar to this one
        for ( similarity , item2 ) in itemMatch [ item ]:

168 # Ignore if this user has already rated this item
            if item2 in userRatings: continue
            # Weighted sum of rating times similarity
            scores.setdefault( item2, 0 )
            scores [ item2 ] += similarity * rating
173 # Sum of all the similarities
            totalSim.setdefault( item2, 0 )
            totalSim [ item2 ] += similarity

178 # Divide each total score by total weighting to get an average
    rankings=[( score / totalSim [ item ] , item ) for item , score in scores.items( )]

    # Return the rankings from highest to lowest
    rankings.sort( )
183 rankings.reverse( )
    return rankings

def loadMovieLens():
    # Get movie titles
188 movies={}
    for line in open( 'u.item' ):
        ( id , title ) = line.split( '|' ) [ 0 : 2 ]
        movies [ id ] = title
    genre = {}
193 for line in open( 'u.item' ):
        ( id , title , release , video , url , unknown , action , adventure , animation , childrens ,
            comedy , crime , documentary , drama , fantasy , filmnoir , horror , musical , mystery ,
            romance , scifi , thriller , war , western ) = line.split( '|' ) [ 0 : 46 ]
        if int( action ) == 1:
            genre [ title ] = 'action'
        elif int( adventure ) == 1:
            genre [ title ] = 'adventure'
198 elif int( animation ) == 1:

```

```

        genre[title]= 'animation'
    elif int(childrens) == 1:
        genre[title]= 'childrens'
203 elif int(comedy) == 1:
        genre[title]= 'comedy'
    elif int(crime) == 1:
        genre[title]= 'crime'
208 elif int(documentary) == 1:
        genre[title]= 'documentary'
    elif int(drama) == 1:
        genre[title]= 'drama'
    elif int(fantasy) == 1:
        genre[title]= 'fantasy'
213 elif int(filmnoir) == 1:
        genre[title]= 'filmnoir'
    elif int(horror) == 1:
        genre[title]= 'horror'
218 elif int(musical) == 1:
        genre[title]= 'musical'
    elif int(romance) == 1:
        genre[title]= 'romance'
    elif int(thriller) == 1:
        genre[title]= 'thriller'
223 elif int(scifi) == 1:
        genre[title]= 'scifi'
    elif int(western) == 1:
        genre[title]= 'western'
228 elif int(war) == 1:
        genre[title]= 'war'
# Load data
prefs={}
for line in open('u.data'):
    (user, movieid, rating, ts)=line.split('\t')
233 prefs.setdefault(user, {})
    prefs[user][movies[movieid]]=float(rating)

users={}
for line in open('u.user'):
238 (user, age, gender, job, zipcode) = line.split('|')
    users.setdefault(user, {})
    users[user] = {'age': age, 'gender': gender, 'job': job, 'zipcode': zipcode}

mfavs = {}
lfavs={}
243 for line in open('u.item'):
    (id, title)= line.split('|')[0:2]
    if str(title) == 'Godfather, The (1972)':
        mfavs.setdefault(title)
248 elif str(title) == 'Shawshank Redemption, The (1994)':
        lfavs.setdefault(title)
return prefs, movies, genre, users, mfavs, lfavs

def get_avg(prefs, mid, user_filter=lambda x: True):
253 ratings = []
    for user, user_ratings in prefs.iteritems():
        if user_filter(users[user]) and user_ratings.has_key(movies[mid]):
            ratings.append(user_ratings[movies[mid]])
    if not ratings:
258 return 0.0
    return mean(ratings)

def get_top(sorted_list, key=lambda x, i: x[i][1], n=5):
263 top = key(sorted_list, 0)
    top_items = []
    i = 0
    while i < n or key(sorted_list, i) == top:
        top_items.append(sorted_list[i])

```

```

268         if i < n and key(sorted_list, i) != top:
                top = key(sorted_list, i)
            i += 1
        return top_items

273 def count_movie_ratings(prefs, mid, transform=False):
    num = 0
    for user, user_ratings in prefs.iteritems():
        if user_ratings.has_key(movies[mid]):
            num += 1
278     return num

def get_sim_ratings(title, similar, top_key=lambda x, i: x[i][1], n=2000):
    itemPrefs = transformPrefs(prefs)
    matches = topMatches(itemPrefs, title, n=n, similarity=sim_pearson)
283     sorted_m = sorted(matches, key=itemgetter(0), reverse=similar)
    return get_top(sorted_m, key=top_key, n=20)

def tabulate(tuples, caption, label, colnames, output):
    output.write('\n\\begin{table}[h!]\n')
    output.write('\n\\centering\n')
    opts = '| ' + ' | '.join(['l' for i in xrange(len(tuples[0]))]) + ' | '
    output.write('\n\\begin{tabular}{{{0}}}\n'.format(opts))
    output.write('\n\\hline\n')
    header = '& '.join(['{}' for i in xrange(len(tuples[0]))]).format(*colnames)
293     output.write(header + '\n\\hline\n')
    for item in tuples:
        temp = '& '.join(['{}' for i in xrange(len(item))])
        output.write(temp.format(*item) + '\n\\hline\n')
    output.write('\n\\hline\n\\end{tabular}\n')
298     output.write('\n\\caption{{{0}}}\n'.format(caption))
    output.write('\n\\label{tab:{0}}\n'.format(label))
    output.write('\n\\end{table}\n\n')

def movies_notRated(prefs, movies, userid, movie_type):
303     userRatings = prefs[userid]
    count = 0
    movie_rated = {}
    for id, title in movies.items():
        temp = 0
308         for item, rating in userRatings.items():
            if title == item:
                temp = 1
                movie_rated.setdefault(title, rating)
        for name, genre in movie_type.items():
313             if title == name:
                if temp == 0:
                    try:
                        if genre == 'comedy':
318                             movie_rated[title] = '3.0'
                            temp = 0
                        elif genre == 'action':
                            movie_rated[title] = '3.0'
                            temp = 0
                        elif genre == 'crime':
323                             movie_rated[title] = '3.0'
                            temp = 0
                        elif genre == 'drama':
                            movie_rated[title] = '4.0'
                            temp = 0
                        elif genre == 'thriller':
328                             movie_rated[title] = '5.0'
                            temp = 0
                        elif genre == 'scifi':
                            movie_rated[title] = '4.0'
                            temp = 0
                        elif genre == 'war':
333                             movie_rated[title] = '3.0'

```

```

338         temp = 0
        elif genre == 'mystery':
            movie_rated[title] = '3.0'
            temp = 0
        else:
            movie_rated[title] = '1.0'
            temp = 0
343     except:
        pass
    return movie_rated

def pearson(avgl, avg2, prefs, movie1, movie2, movies):
348     r={}
    r=0
    top = 0
    rbtm = 0
    lbtm = 0
353
    for user, movie in prefs.items():
        userRatings = prefs[user]
        for item, rating in userRatings.items():
            if movie1 == movies[item]:
358                 for user1, move in prefs.items():
                    userRatings = prefs[user1]
                    for item1, rating1 in userRatings.items():
                        if movie2 == movies[item1]:
                            if user1 == user:
363                                 top += (rating-avgl)*(rating1-avg2)
                                    lbtm += (rating-avgl)*(rating1-avgl)
                                    rbtm += (rating1-avg2)*(rating1-avg2)

    if lbtm == 0 or rbtm == 0:
368         r = 5
        return r
    else:
        r = top/(math.sqrt(lbtm)* math.sqrt(rbtm))
    return r
373

print "Parsing data"
prefs, movies, genre, users, mfavs, lfavs = loadMovieLens()

def flatten(tup, f=lambda x: (x[0], x[1][0][1], x[1][0][0])):
378     return f(tup)

if __name__ == '__main__':
    with open('output.tex', 'w') as outfile:
        getuser = {}
        user_filter = lambda x: x['gender'] == 'M' and x['job'] == 'student' and int(x['age'
383            ]) == 23
        ratings = []
        for mid, movie in movies.iteritems():
            for user, user_ratings in prefs.iteritems():
                if user_filter(users[user]) and user_ratings.has_key(movies[mid]):
388                     getuser.setdefault(int(user), {})
                        getuser[int(user)][movies[mid]] = float(user_ratings[movies[mid]])
                        ratings.append(user_ratings[movies[mid]])

        sorted_getuser = {}
        movie_sort = {}
393     for user, user_movie in getuser.items():
        user_movie_sort = sorted(user_movie.items(), key=itemgetter(1), reverse=False)
        for title, rating in user_movie_sort[:3]:
            movie_sort.setdefault(title, rating)

        # print movie_sort
398     sorted_getuser.setdefault(user, user_movie_sort[:3])
    sorted_avg_all = sorted(sorted_getuser.items(), key=itemgetter(0), reverse=False)
    # print sorted_avg_all
    top_raters = get_top(sorted_avg_all, key=lambda x,i: x[i][1][0])
    top_raters = [flatten(rater) for rater in sorted_avg_all]

```

```

403 tabulate(top_raters, 'Users', 'user', ('User', 'Rating', 'Movie'), outfile)
    print "done with 1"
    most_correlated = {}
    least_correlated = {}
    for user, rest in users.iteritems():
408         if int(user) == 135:
            for user1, rest in users.iteritems():
                if user == user1:
                    pass
                else:
413                     r = sim_pearson(prefs, user, user1)
                        if r == 1.0:
                            most_correlated.setdefault(int(user1), r)
                        if r == -1.0:
                            least_correlated.setdefault(int(user1), r)
418 sorted_most_correlated = sorted(most_correlated.items(), key=itemgetter(0), reverse=
    False)[len(most_correlated) - 6:-1]
    sorted_least_correlated = sorted(least_correlated.items(), key=itemgetter(0),
    reverse=False)[len(least_correlated) - 6:-1]
    tabulate(sorted_most_correlated, 'Most Correlated', 'most', ('User', 'Pearson\'s r')
    , outfile)
    tabulate(sorted_least_correlated, 'Least Correlated', 'least', ('User', 'Pearson\'s
    r'), outfile)
    print "done with 2"
    movieRated1 = {}
423 for user, rest in users.iteritems():
        if int(user) == 135:
            movieRated1 = movies_not_rated(prefs, movies, user, genre)
            sorted_most_recomm = sorted(movieRated1.items(), key=itemgetter(1), reverse=False)[
            len(movieRated1) - 6:-1]
428 sorted_least_recomm = sorted(movieRated1.items(), key=itemgetter(1), reverse=True)[
            len(movieRated1) - 6:-1]
            tabulate(sorted_most_recomm, 'Top 5 unseen movies recommendations', 'recommtop', (
            'Title', 'Recomm Ranking'), outfile)
            tabulate(sorted_least_recomm, 'Least 5 unseen movies recommendations', 'recommleast'
            , ('Title', 'Recomm Ranking'), outfile)
            print "done with 3"
            fav_cor = {}
            nfav_cor = {}
433 avg1 = get_avg(prefs, mfavs)
            for id, title in movies.items():
                avg2 = get_avg(prefs, id)
                r = pearson(avg1, avg2, prefs, mfavs, title)
438                 if r == 1.0:
                    fav_cor.setdefault(str(title), str(r))
                elif r == -1.0:
                    nfav_cor.setdefault(str(title), str(r))
            sorted_fav_correlated = sorted(fav_cor.items(), key=itemgetter(0), reverse=False)[
            len(fav_cor) - 6:-1]
443 sorted_nfav_correlated = sorted(nfav_cor.items(), key=itemgetter(0), reverse=False)[
            len(nfav_cor) - 6:-1]
            tabulate(sorted_fav_correlated, 'Most Favourite Movie \textcolor{blue}{Godfather,
            The (1972)} Correlated', 'fav', ('Title', 'Pearson\'s r'), outfile)
            tabulate(sorted_nfav_correlated, 'Least Favourite Movie \textcolor{blue}{Godfather,
            The (1972)} Correlated', 'nfav', ('Title', 'Pearson\'s r'), outfile)
            print "done with 4"

```

Listing 12: all code used

## References

- [1] Numpy Developers. Python numpy module. <http://www.numpy.org/>, 2016.