

CS532 Web Science: Assignment 8

Finished on April 7, 2016

Dr. Michael L. Nelson

Naina Sai Tipparti
ntippart@cs.odu.edu

Contents

| | |
|--------------------|----------|
| Problem 1 | 2 |
| Question | 2 |
| Answer | 2 |
| Problem 2 | 6 |
| Question | 6 |
| Answer | 6 |
| Appendix A | 9 |

Listings

| | | |
|----|-----------------------------------------------|----|
| 1 | referenced variables in get_uris.py | 2 |
| 2 | Sample list of Blog URI's | 2 |
| 3 | main for get_uris.py | 3 |
| 4 | get_atom function | 3 |
| 5 | add_uri function | 3 |
| 6 | looping over the URIs | 3 |
| 7 | processing each blog | 4 |
| 8 | creating the blog data matrix | 4 |
| 9 | loading the data | 5 |
| 10 | building the master wordlist | 5 |
| 11 | writing the data | 5 |
| 12 | creating the dendrograms | 6 |
| 13 | creating the dendrograms | 6 |
| 14 | hcluster function | 6 |
| 15 | printclust function | 7 |
| 16 | drawdendrogram function | 7 |
| 17 | get_uris.py | 9 |
| 18 | matrix.py | 10 |
| 19 | clusters.py | 12 |

List of Figures

| | | |
|---|---------------------------|---|
| 1 | blog dendrogram | 8 |
|---|---------------------------|---|

Problem 1

Question

Create a blog-term matrix. Start by grabbing 100 blogs; include:

<http://f-measure.blogspot.com/>
<http://ws-dl.blogspot.com/>

and grab 98 more as per the method shown in class. Note that this method randomly chooses blogs and each student will separately do this process, so it is unlikely that these 98 blogs will be shared among students. In other words, no sharing of blog data. Upload to github your code for grabbing the blogs and provide a list of blog URIs, both in the report and in github..

Use the blog title as the identifier for each blog (and row of the matrix). Use the terms from every item/title (RSS) or entry/title (Atom) for the columns of the matrix. The values are the frequency of occurrence. Essentially you are replicating the format of the “blog-data.txt” file included with the PCI book code. Limit the number of terms to the most “popular” (i.e., frequent) 500 terms, this is *after* the criteria on p. 32 (slide 7) has been satisfied.

Answer

To complete this assignment, a blog word count matrix was required. To start off, a list of blog URIs was obtained using the method described in class, implemented as the `get_uris.py` script, which can be found in Appendix A, Listing 17. Two default blogs, F-Measure and the Old Dominion Web Science and Digital Libraries blogs, were added as defaults to the initial URI list and then, using the seed URI provided (Listing 1), the remaining 98 URIs from random blogs within the blogger.com family were added.

```
default = 'http://www.blogger.com/next-blog?navBar=true&blogID=3471633091411211117'
must_haves = ['http://f-measure.blogspot.com/', 'http://ws-dl.blogspot.com/']
```

Listing 1: referenced variables in `get_uris.py`

The `get_uris` main function in Listing 3 was the driver that called the `get_atom` function (shown in Listing 4) to extract the atom [1] URIs from each blog and add them to the set of URIs with the `add_uri` function, shown in Listing 5.

```
20 http://jbalow.blogspot.com/feeds/posts/default
   http://toxicreine.blogspot.com/feeds/posts/default
   http://beautifulsweetpea.blogspot.com/feeds/posts/default
   http://brittanyvsutah.blogspot.com/feeds/posts/default
   http://thewoolfpack6.blogspot.com/feeds/posts/default
   http://jennie-tabs-wedding.blogspot.com/feeds/posts/default
25 http://kamielverwer.blogspot.com/feeds/posts/default
   http://azstampcrazy.blogspot.com/feeds/posts/default
   http://deneiserothenberger.blogspot.com/feeds/posts/default
   http://gobobgo-anewday.blogspot.com/feeds/posts/default
```

Listing 2: Sample list of Blog URI's

```

if __name__ == '__main__':
    uris = set()
    with open('blog_uris', 'a') as outfile:
        if len(sys.argv) > 1 and sys.argv[1] == 'new':
            for must_have in must_haves:
                uri = get_atom(must_have)
                add_uri(uri, uris, outfile)
        else:
            with open('blog_uris') as infile:
                [uris.add(line.strip()) for line in infile]
            while len(uris) < 100:
                uri = get_atom(default)
                add_uri(uri, uris, outfile)

```

Listing 3: main for get_uris.py

```

10 def get_atom(uri):
    try:
        r = requests.get(uri)
    except Exception, e:
        return None
15 soup = BeautifulSoup(r.text)
    links = soup.find_all('link', {'type': 'application/atom+xml'})
    if links:
        return str(links[0]['href'])
    return None

```

Listing 4: get_atom function

```

def add_uri(uri, uris, outfile):
    if uri and uri not in uris:
        uris.add(uri)
        outfile.write(uri + '\n')
25 print len(uris), uri

```

Listing 5: add_uri function

After the full list of 100 URIs was obtained, page counts for each blog were extracted and saved to a file called `pagecounts` using the `matrix.py` script. This script is a modified version of `generatefeedvectors.py` from the book *Programming Collective Intelligence* [2] and can be found in full in Appendix A, Listing 18.

The code responsible for downloading the blogs and counting the words in each is shown in Listing 6, which calls the `get_titles`, `get_words` and `get_next` functions found in Listing 7. This code loops over the list of URIs that was obtained with the `get_uris.py` script (Listing 17), parses each entry, and extracts all the words in each entry's title. These word counts are then saved as a python dictionary to the hard drive for later use.

```

95 if __name__ == '__main__':
    with open('blog_uris') as infile:
        uris = [line.strip() for line in infile if line.strip()]
        if len(sys.argv) == 2 and sys.argv[1] == 'get':
            with futures.ThreadPoolExecutor(max_workers=8) as executor:
                uri_futures = [executor.submit(get_titles, uri) for uri in uris]
                for future in futures.as_completed(uri_futures):
                    uri, title, subtitle, pages, wc = future.result()
                    with open('wcs/' + md5.new(uri).hexdigest(), 'w') as out:
                        out.write(title + ': ' + subtitle + '\t' + str(pages) + '\t')
105 json.dump(wc, out)

```

Listing 6: looping over the URIs

```

def get_next(d):
    for item in d.feed.links:
        if item['rel'] == u'next':
            return item['href']
    return None

def get_words(text):
    txt = re.compile(r'<[>]+>').sub(' ', text)
    words = re.compile(r'^A-Z^a-z+').split(txt)
    return [word.lower() for word in words if word != '']

def get_titles(uri):
    print('processing {}'.format(uri))
    next = uri
    wc = {}
    pages = 0
    while next is not None:
        d = feedparser.parse(next)
        for e in d.entries:
            words = get_words(e.title.encode('utf-8'))
            for word in words:
                wc.setdefault(word, 0)
                wc[word] += 1
        pages += 1
        next = get_next(d)
        print('next {}'.format(next))
    title = d.feed.title.encode('utf-8')
    subtitle = d.feed.subtitle[:50].encode('utf-8')
    print('finished: {}: {}'.format(title, subtitle))
    return uri, title, subtitle, pages, wc

```

Listing 7: processing each blog

The parsed results were then read by the code in Listing 8. This code used the `load_data` and `build_wordlist` functions in Listing 9 and 10 to read each of the blog word counts and then created four collections to organize them all:

1. **apcount**: A dictionary containing the count for all words combined
2. **wordcounts**: A dictionary containing each blog's individual word count
3. **pagecounts**: A dictionary containing each blog's page count
4. **wordlist**: A list containing all of the words found in each blog

```

apcount, wordcounts, pagecounts = load_data(uris)
wordlist = build_wordlist(apcount, uris)
if len(sys.argv) == 2 and sys.argv[1] == 'pages':
    with open('pagecounts', 'w') as outfile:
        outfile.write('blog\tpages\n')
        for blog, pagecount in pagecounts.iteritems():
            outfile.write("{}{} + blog.replace("\\", "") + "\\t" + str(
                pagecount) + "\\n")
elif len(sys.argv) == 2 and sys.argv[1] == 'wc':
    write_data('blogdata1.txt', wordlist, wordcounts)

```

Listing 8: creating the blog data matrix

```

50 def load_data(uris):
    apcount = {}
    wordcounts = {}
    pagecounts = {}
    for uri in uris:
55         with open('wcs/' + md5.new(uri).hexdigest()) as infile:
            try:
                lines = infile.read().split('\t')
                title = lines[0]
                pages = int(lines[1])
60                 wc = json.loads(lines[2])
            except Exception, e:
                print('*** {} generated an exception: {}'.format(uri, e))
                continue
            wordcounts[title] = wc
65             pagecounts[title] = pages
            for word, count in wc.items():
                apcount.setdefault(word, 0)
                apcount[word] += count
    return apcount, wordcounts, pagecounts

```

Listing 9: loading the data

```

def build_wordlist(apcount, uris):
    wordlist = []
    for w, bc in sorted(apcount.items(), key=lambda x: x[1], reverse=True):
75         frac = float(bc) / len(uris)
        if frac > 0.1 and frac < 0.5:
            wordlist.append(w)
    return wordlist

```

Listing 10: building the master wordlist

The code in Listing 11 then created the matrix using the `write_data` function using the data structures that store the blog word counts.

```

80 def write_data(filename, wordlist, wordcounts, form=lambda wc, word, wordcounts: wc[word]):
    with open(filename, 'w') as out:
        out.write('Blog')
        for word in wordlist[:500]:
            out.write('\t%s' % word)
            out.write('\n')
85         for blog, wc in wordcounts.items():
            print blog
            out.write(blog)
            for word in wordlist[:500]:
                if word in wc:
90                     out.write('\t{}'.format(form(wc, word, wordcounts)))
                else:
                    out.write('\t0')
            out.write('\n')

```

Listing 11: writing the data

Problem 2

Question

Create an ASCII and JPEG dendrogram that clusters (i.e., HAC) the most similar blogs (see slides 12 & 13). Include the JPEG in your report and upload the ascii file to github (it will be too unwieldy for inclusion in the report).

Answer

The ascii and jpeg dendrograms were created using the code shown in Listing 12, which is modeled after the example from class.

```

290 blognames, words, data = readfile('q1/blogdata1.txt')
    clust = hcluster(data)
    with open('dendrogram.txt', 'w') as outfile:
        stdout = sys.stdout
        sys.stdout = outfile
        printclust(clust, labels=blognames)
        sys.stdout = stdout
    drawdendrogram(clust, blognames, jpeg='blogclust.jpg')
```

Listing 12: creating the dendrograms

The `readfile` function shown in Listing 13 was used to read the data that was compiled from Question 1 into memory where it is then processed by the `hcluster` function found in Listing 14 to produce the clustered representation of the blogs.

```

5 def readfile(filename):
    lines=[line for line in file(filename)]

    # First line is the column titles
    colnames=lines[0].strip().split('\t')[1:]
    rownames=[]
    data=[]
10 for line in lines[1:]:
    p=line.strip().split('\t')
    # First column in each row is the rowname
    rownames.append(p[0])
    # The data for this row is the remainder of the row
15 data.append([float(x) for x in p[1:]])
    return rownames, colnames, data
```

Listing 13: creating the dendrograms

```

50 def hcluster(rows, distance=pearson):
    distances={}
    currentclustid=-1

    # Clusters are initially just the rows
    clust=[bicluster(rows[i], id=i) for i in range(len(rows))]

55 while len(clust)>1:
    lowestpair=(0,1)
    closest=distance(clust[0].vec, clust[1].vec)

    # loop through every pair looking for the smallest distance
60 for i in range(len(clust)):
    for j in range(i+1, len(clust)):
        # distances is the cache of distance calculations
        if (clust[i].id, clust[j].id) not in distances:
```

```

        distances[(clust[i].id, clust[j].id)] = distance(clust[i].vec, clust[j].vec)
65
        d = distances[(clust[i].id, clust[j].id)]

        if d < closest:
            closest = d
70
            lowestpair = (i, j)

        # calculate the average of the two clusters
        mergevec = [
            (clust[lowestpair[0]].vec[i] + clust[lowestpair[1]].vec[i]) / 2.0
75
        for i in range(len(clust[0].vec))]

        # create the new cluster
        newcluster = bicluster(mergevec, left=clust[lowestpair[0]],
80
                               right=clust[lowestpair[1]],
                               distance=closest, id=currentclustid)

        # cluster ids that weren't in the original set are negative
        currentclustid -= 1
        del clust[lowestpair[1]]
85
        del clust[lowestpair[0]]
        clust.append(newcluster)

    return clust[0]

```

Listing 14: hcluster function

The `printclust` function from Listing 15 prints the ascii dendrogram of the cluster object parameter to `sys.stdout`, which is redirected to write to a file with the code in Listing 12.

```

90 def printclust(clust, labels=None, n=0):
    # indent to make a hierarchy layout
    for i in range(n): print ' ',
    if clust.id < 0:
        # negative id means that this is branch
95
        print '-'
    else:
        # positive id means that this is an endpoint
        if labels == None: print clust.id
        else: print labels[clust.id]
100

    # now print the right and left branches
    if clust.left != None: printclust(clust.left, labels=labels, n=n+1)
    if clust.right != None: printclust(clust.right, labels=labels, n=n+1)

```

Listing 15: printclust function

The `drawdendrogram` function from Listing 16 creates a jpeg image of the cluster, which is shown in Figure 1.

```

def drawdendrogram(clust, labels, jpeg='clusters.jpg'):
    # height and width
    h = getheight(clust) * 20
125
    w = 1200
    depth = getdepth(clust)

    # width is fixed, so scale distances accordingly
    scaling = float(w - 150) / depth
130

    # Create a new image with a white background
    img = Image.new('RGB', (w, h), (255, 255, 255))
    draw = ImageDraw.Draw(img)

135
    draw.line((0, h/2, 10, h/2), fill=(255, 0, 0))

```


Listing 16: drawdendrogram function



Appendix A

```
1  #!/usr/bin/env python

import requests
import sys
from bs4 import BeautifulSoup

6  default = 'http://www.blogger.com/next-blog?navBar=true&blogID=3471633091411211117'
must_haves = ['http://f-measure.blogspot.com/', 'http://ws-dl.blogspot.com/']

def get_atom(uri):
11  try:
        r = requests.get(uri)
    except Exception, e:
        return None
    soup = BeautifulSoup(r.text)
16  links = soup.find_all('link', {'type': 'application/atom+xml'})
    if links:
        return str(links[0]['href'])
    return None

21 def add_uri(uri, uris, outfile):
    if uri and uri not in uris:
        uris.add(uri)
        outfile.write(uri + '\n')
        print len(uris), uri

26 if __name__ == '__main__':
    uris = set()
    with open('blog_uris', 'a') as outfile:
        if len(sys.argv) > 1 and sys.argv[1] == 'new':
31         for must_have in must_haves:
            uri = get_atom(must_have)
            add_uri(uri, uris, outfile)
        else:
            with open('blog_uris') as infile:
36             [uris.add(line.strip()) for line in infile]
    while len(uris) < 100:
        uri = get_atom(default)
        add_uri(uri, uris, outfile)
```

Listing 17: get_uris.py

```

1 import feedparser
import futures
import math
import md5
import re
6 import sys
import json

def get_next(d):
    for item in d.feed.links:
11         if item['rel'] == u'next':
            return item['href']
    return None

def get_words(text):
16     txt = re.compile(r'<[>]+>').sub(' ', text)
    words = re.compile(r'^A-Za-z+').split(txt)
    return [word.lower() for word in words if word != '']

def get_titles(uri):
21     print('processing {}'.format(uri))
    next = uri
    wc = {}
    pages = 0
    while next is not None:
26         d = feedparser.parse(next)
        for e in d.entries:
            words = get_words(e.title.encode('utf-8'))
            for word in words:
                wc.setdefault(word, 0)
31                 wc[word] += 1
            pages += 1
            next = get_next(d)
            print('next {}'.format(next))
        title = d.feed.title.encode('utf-8')
        subtitle = d.feed.subtitle[:50].encode('utf-8')
36         print('finished: {}: {}'.format(title, subtitle))
    return uri, title, subtitle, pages, wc

def tf(wc, word):
41     return float(wc[word]) / float(sum(wc.values()))

def idf(wordcounts, word):
    present = 0
    for wc in wordcounts.values():
46         if word in wc:
            present += 1
    return math.log(len(wordcounts) / present, 2)

def load_data(uris):
51     apcount = {}
    wordcounts = {}
    pagecounts = {}
    for uri in uris:
        with open('wcs/' + md5.new(uri).hexdigest()) as infile:
56             try:
                lines = infile.read().split('\n')
                title = lines[0]
                pages = int(lines[1])
                wc = json.loads(lines[2])
61             except Exception, e:
                print('*** {} generated an exception: {}'.format(uri, e))
                continue
            wordcounts[title] = wc
            pagecounts[title] = pages
66             for word, count in wc.items():
                apcount.setdefault(word, 0)

```

```

        apcount[word] += count
    return apcount, wordcounts, pagecounts

71 def build_wordlist(apcount, uris):
    wordlist = []
    for w, bc in sorted(apcount.items(), key=lambda x: x[1], reverse=True):
        frac = float(bc) / len(uris)
        if frac > 0.1 and frac < 0.5:
76         wordlist.append(w)
    return wordlist

def write_data(filename, wordlist, wordcounts, form=lambda wc, word, wordcounts: wc[word]):
    with open(filename, 'w') as out:
81         out.write('Blog')
        for word in wordlist[:500]:
            out.write('\t%s' % word)
        out.write('\n')
        for blog, wc in wordcounts.items():
86             print blog
            out.write(blog)
            for word in wordlist[:500]:
                if word in wc:
                    out.write('\t{}'.format(form(wc, word, wordcounts)))
91                 else:
                    out.write('\t0')
            out.write('\n')

if __name__ == '__main__':
96     with open('blog_uris') as infile:
        uris = [line.strip() for line in infile if line.strip()]
        if len(sys.argv) == 2 and sys.argv[1] == 'get':
            with futures.ThreadPoolExecutor(max_workers=8) as executor:
                uri_futures = [executor.submit(get_titles, uri) for uri in uris]
101             for future in futures.as_completed(uri_futures):
                uri, title, subtitle, pages, wc = future.result()
                with open('wcs/' + md5.new(uri).hexdigest(), 'w') as out:
                    out.write(title + ': ' + subtitle + '\t' + str(pages) + '\t')
                    json.dump(wc, out)
106         else:
            apcount, wordcounts, pagecounts = load_data(uris)
            wordlist = build_wordlist(apcount, uris)
            if len(sys.argv) == 2 and sys.argv[1] == 'pages':
                with open('pagecounts', 'w') as outfile:
111                     outfile.write('blog\tpages\n')
                    for blog, pagecount in pagecounts.iteritems():
                        outfile.write("\"" + blog.replace("\", \"") + "\"" + "\t" + str(
                            pagecount) + '\n')
            elif len(sys.argv) == 2 and sys.argv[1] == 'wc':
                write_data('blogdata1.txt', wordlist, wordcounts)
116             elif len(sys.argv) == 2 and sys.argv[1] == 'tfidf':
                write_data('blogdata2.txt', wordlist, wordcounts, form=lambda wc, word,
                    wordcounts: tf(wc, word) * idf(wordcounts, word))

```

Listing 18: matrix.py

```

from PIL import Image, ImageDraw

3 def readfile(filename):
    lines=[line for line in file(filename)]

    # First line is the column titles
    colnames=lines[0].strip().split('\t')[1:]
    8 rownames=[]
    data=[]
    for line in lines[1:]:
        p=line.strip().split('\t')
        # First column in each row is the rowname
        13 rownames.append(p[0])
        # The data for this row is the remainder of the row
        data.append([float(x) for x in p[1:]])
    return rownames, colnames, data

18
from math import sqrt

def pearson(v1,v2):
    # Simple sums
    23 sum1=sum(v1)
    sum2=sum(v2)

    # Sums of the squares
    sum1Sq=sum([pow(v,2) for v in v1])
    28 sum2Sq=sum([pow(v,2) for v in v2])

    # Sum of the products
    pSum=sum([v1[i]*v2[i] for i in range(len(v1))])

    33 # Calculate r (Pearson score)
    num=pSum-(sum1*sum2/len(v1))
    den=sqrt((sum1Sq-pow(sum1,2)/len(v1))*(sum2Sq-pow(sum2,2)/len(v1)))
    if den==0: return 0

    38 return 1.0-num/den

class bicluster:
    def __init__(self, vec, left=None, right=None, distance=0.0, id=None):
        43 self.left=left
        self.right=right
        self.vec=vec
        self.id=id
        self.distance=distance

    48 def hcluster(rows, distance=pearson):
        distances={}
        currentclustid=-1

        # Clusters are initially just the rows
        53 clust=[bicluster(rows[i], id=i) for i in range(len(rows))]

        while len(clust)>1:
            lowestpair=(0,1)
            closest=distance(clust[0].vec, clust[1].vec)

            58 # loop through every pair looking for the smallest distance
            for i in range(len(clust)):
                for j in range(i+1, len(clust)):
                    # distances is the cache of distance calculations
                    63 if (clust[i].id, clust[j].id) not in distances:
                        distances[(clust[i].id, clust[j].id)]=distance(clust[i].vec, clust[j].vec)

                    d=distances[(clust[i].id, clust[j].id)]

```

```

68         if d<closest:
            closest=d
            lowestpair=(i,j)

        # calculate the average of the two clusters
73         mergevec=[
            (clust[lowestpair[0]].vec[i]+clust[lowestpair[1]].vec[i])/2.0
            for i in range(len(clust[0].vec))]

        # create the new cluster
78         newcluster=biclust(mergevec, left=clust[lowestpair[0]],
                               right=clust[lowestpair[1]],
                               distance=closest, id=currentclustid)

        # cluster ids that weren't in the original set are negative
83         currentclustid -=1
        del clust[lowestpair[1]]
        del clust[lowestpair[0]]
        clust.append(newcluster)

88     return clust[0]

def printclust(clust, labels=None, n=0):
    # indent to make a hierarchy layout
    for i in range(n): print ' ',
93     if clust.id<0:
        # negative id means that this is branch
        print '└─'
    else:
        # positive id means that this is an endpoint
98         if labels==None: print clust.id
        else: print labels[clust.id]

    # now print the right and left branches
    if clust.left!=None: printclust(clust.left, labels=labels, n=n+1)
103    if clust.right!=None: printclust(clust.right, labels=labels, n=n+1)

def getheight(clust):
    # Is this an endpoint? Then the height is just 1
    if clust.left==None and clust.right==None: return 1
108
    # Otherwise the height is the same of the heights of
    # each branch
    return getheight(clust.left)+getheight(clust.right)

113 def getdepth(clust):
    # The distance of an endpoint is 0.0
    if clust.left==None and clust.right==None: return 0

    # The distance of a branch is the greater of its two sides
118    # plus its own distance
    return max(getdepth(clust.left), getdepth(clust.right))+clust.distance

def drawdendrogram(clust, labels, jpeg='clusters.jpg'):
123    # height and width
    h=getheight(clust)*20
    w=1200
    depth=getdepth(clust)

128    # width is fixed, so scale distances accordingly
    scaling=float(w-150)/depth

    # Create a new image with a white background
    img=Image.new('RGB', (w,h), (255,255,255))
133    draw=ImageDraw.Draw(img)

    draw.line((0, h/2, 10, h/2), fill=(255,0,0))

```

```

138 # Draw the first node
drawnode(draw, clust, 10, (h/2), scaling, labels)
img.save(jpeg, 'JPEG')

def drawnode(draw, clust, x, y, scaling, labels):
    if clust.id < 0:
143         h1=getheight(clust.left)*20
         h2=getheight(clust.right)*20
         top=y-(h1+h2)/2
         bottom=y+(h1+h2)/2
         # Line length
148         ll=clust.distance*scaling
         # Vertical line from this cluster to children
         draw.line((x, top+h1/2, x, bottom-h2/2), fill=(255,0,0))

         # Horizontal line to left item
153         draw.line((x, top+h1/2, x+ll, top+h1/2), fill=(255,0,0))

         # Horizontal line to right item
         draw.line((x, bottom-h2/2, x+ll, bottom-h2/2), fill=(255,0,0))

158         # Call the function to draw the left and right nodes
         drawnode(draw, clust.left, x+ll, top+h1/2, scaling, labels)
         drawnode(draw, clust.right, x+ll, bottom-h2/2, scaling, labels)
     else:
         # If this is an endpoint, draw the item label
163         draw.text((x+5, y-7), labels[clust.id], (0,0,0))

def rotatematrix(data):
    newdata=[]
    for i in range(len(data[0])):
168         newrow=[data[j][i] for j in range(len(data))]
         newdata.append(newrow)
    return newdata

import random

173 def kcluster(rows, distance=pearson, k=4):
    # Determine the minimum and maximum values for each point
    ranges=[(min([row[i] for row in rows]), max([row[i] for row in rows]))
    for i in range(len(rows[0]))]

178 # Create k randomly placed centroids
    clusters=[[random.random()*(ranges[i][1]-ranges[i][0])+ranges[i][0]
    for i in range(len(rows[0]))] for j in range(k)]

183 lastmatches=None
    for t in range(100):
        print 'Iteration %d' % t
        bestmatches=[] for i in range(k)]

188 # Find which centroid is the closest for each row
        for j in range(len(rows)):
            row=rows[j]
            bestmatch=0
            for i in range(k):
193                 d=distance(clusters[i], row)
                 if d<distance(clusters[bestmatch], row): bestmatch=i
            bestmatches[bestmatch].append(j)

        # If the results are the same as last time, this is complete
198         if bestmatches==lastmatches: break
        lastmatches=bestmatches

        # Move the centroids to the average of their members
        for i in range(k):
203             avgs=[0.0]*len(rows[0])

```

```

    if len(bestmatches[i])>0:
        for rowid in bestmatches[i]:
            for m in range(len(rows[rowid])):
                avgs[m]+=rows[rowid][m]
208         for j in range(len(avgs)):
                avgs[j]/=len(bestmatches[i])
                clusters[i]=avgs

    return bestmatches
213
def tanamoto(v1,v2):
    c1,c2,shr=0,0,0

    for i in range(len(v1)):
218         if v1[i]!=0: c1+=1 # in v1
                if v2[i]!=0: c2+=1 # in v2
                if v1[i]!=0 and v2[i]!=0: shr+=1 # in both

    return 1.0-(float(shr)/(c1+c2-shr))
223
def scaledown(data,distance=pearson,rate=0.01):
    n=len(data)

    # The real distances between every pair of items
228     realdist=[[distance(data[i],data[j]) for j in range(n)]
                for i in range(0,n)]

    # Randomly initialize the starting points of the locations in 2D
    loc=[[random.random(),random.random()] for i in range(n)]
233     fakedist=[[0.0 for j in range(n)] for i in range(n)]

    lasterror=None
    for m in range(0,1000):
        # Find projected distances
238         for i in range(n):
                for j in range(n):
                        fakedist[i][j]=sqrt(sum([pow(loc[i][x]-loc[j][x],2)
                                                for x in range(len(loc[i]))]))

243     # Move points
    grad=[[0.0,0.0] for i in range(n)]

    totalerror=0
    for k in range(n):
248         for j in range(n):
                if j==k: continue
                # The error is percent difference between the distances
                if realdist[j][k] != 0:
                        errorterm=(fakedist[j][k]-realdist[j][k])/realdist[j][k]

253         # Each point needs to be moved away from or towards the other
                # point in proportion to how much error it has
                grad[k][0]+=((loc[k][0]-loc[j][0])/fakedist[j][k])*errorterm
                grad[k][1]+=((loc[k][1]-loc[j][1])/fakedist[j][k])*errorterm

258         # Keep track of the total error
                totalerror+=abs(errorterm)
    print totalerror

263     # If the answer got worse by moving the points, we are done
    if lasterror and lasterror<totalerror: break
    lasterror=totalerror

    # Move each of the points by the learning rate times the gradient
268     for k in range(n):
            loc[k][0]-=rate*grad[k][0]
            loc[k][1]-=rate*grad[k][1]

```



```

273     return loc

def draw2d(data, labels, jpeg='mds2d.jpg'):
    img=Image.new('RGB',(2000,2000),(255,255,255))
    draw=ImageDraw.Draw(img)
    for i in range(len(data)):
278         x=(data[i][0]+0.5)*1000
            y=(data[i][1]+0.5)*1000
            draw.text((x,y),labels[i],(0,0,0))
    img.save(jpeg,'JPEG')

283 import sys

if __name__ == '__main__':
    blognames, words, data = readfile('q1/blogdata1.txt')
    clust = hcluster(data)
288     with open('dendrogram.txt','w') as outfile:
        stdout = sys.stdout
        sys.stdout = outfile
        printclust(clust, labels=blognames)
        sys.stdout = stdout
293     drawdendrogram(clust, blognames, jpeg='blogclust.jpg')
    print "Done with dendrograms"
    outfile=open('kclust.txt','w')
    for i in [5,10,20]:
        kclust, iternum=kcluster(data,k=i)
298         outfile.write('\n\n k = %d'%i)
        outfile.write('Iterations = %d\n'%iternum)
        for cluster in kclust:
            outfile.write('[')
            for blogidx in cluster:
303                 outfile.write(blognames[blogidx]+' , ')
            outfile.write(']\n')
    outfile.close()
    coords=scaledown(data)
    draw2d(coords, blognames, jpeg='blogs2d.jpg')

```

Listing 19: clusters.py

References

- [1] Internet Engineering Task Force (IETF). Rfc-4287 the atom syndication format. <https://tools.ietf.org/html/rfc4287>, 2016.
- [2] Toby Segaran. Programming collective intelligence. oăŽreilly, first edition,, 2007.